

50004 - Operating Systems - Lecture 5

Oliver Killane

02/11/21

Process Synchronisation

The key concepts to consider in process synchronisation are:

- **Critical Sections**
Section of code where processes access a shared resource
- **Mutual Exclusion**
Multiple threads/processes cannot execute in a critical section simultaneously.
- **Atomic Operations**
Operations that can occur without interruption or interleave by other threads, e.g reading and writing in 1 instruction.
- **Race Conditions**
Where the behaviour of the program is dependent on the timing and interleaving of threads.
- **Deadlock**
Where the completion of tasks are mutually dependent, there is some loop in the dependency graph meaning a task's completion cannot occur before its completion.
- **Starvation**
- **Synchronisation Mechanisms**
Locks, Semaphores and Monitors for example.

Primitive that allow programmers to enforce ordering of tasks, and exclusion. They are required at entry to and exit from a critical section.

Requirements of Mutual Exclusion

- No two processes can simultaneously be in a critical section.
- No process running outside the critical section can prevent another from entering the critical section.
- When no process is in the critical region, any process requesting permission to enter must be immediately admitted.
- No process requiring access to a critical section can be delayed indefinitely.
- No assumptions are made about the relative speed of processes.

Disabling Interrupts

```
1 void Extract(int acc_no , int sum)
2 {
3     CLI();
4     int B = Acc[acc_no];
5     Acc[acc_no] = B - sum;
6     STI();
7 }
```

- Works only on single-core Systems (other core may access shared in memory).
- May never release the CPU if code buggy.
- Interrupts could be missed while interrupts disabled.
- Only possible in kernel code which has the ability to disable interrupts.

Software Solution - Strict Alternative

<pre>1 while (true) { 2 while (turn != 0) 3 /* loop */ ; 4 critical_section(); 5 turn = 1; 6 noncritical_section0(); 7 }</pre>	<pre>1 while (true) { 2 while (turn != 1) 3 /* loop */ ; 4 critical_section(); 5 turn = 0; 6 noncritical_section0(); 7 }</pre>
--	--

Issues:

- Cannot run 0 or 1 twice in a row without the other working.
- Busy waiting of thread on turn variable.
- If the noncritical section takes a long time, the other thread is blocked as it must wait for turn to be set. (Thread outside critical region prevent critical region access).
- Turn must be volatile to prevent compiler optimisations causing issues.

Busy Waiting

Continuously checking a value to determine if a thread can progress. It is a waste of CPU time and should only be used if the wait is short.

However as a result of actively polling the a value, it can stop waiting very quickly. Spinlocks used in OS kernels take advantage of this.

Peterson's Solution

<pre> 1 int turn = 0; 2 int interested[2] = {0, 0}; 3 4 // thread is 0 or 1 5 void enter_critical(int thread) { 6 int other = 1 - thread; 7 interested[thread] = 1; 8 turn = other; 9 while (turn == other && interested[other]) 10 /* loop */ ; 11 } 12 13 void leave_critical(int thread) { 14 interested[thread] = 0; 15 }</pre>	<pre> 1 enter_critical(1); 2 critical_section(); 3 leave_critical(1);</pre>
	<pre> 1 enter_critical(0); 2 critical_section(); 3 leave_critical(0);</pre>

This lock busy waits while the other is interested and it is the other's turn.

Atomic Operations

```

1 void Extract(int acc_no, int sum)
2 {
3     int B = Acc[acc_no];
4     Acc[acc_no] = B - sum;
5 }
```

Not atomic, other extract could be interleaved.



```

1 void Extract(int acc_no, int sum)
2 {
3     Acc[acc_no] -= sum;
4 }
```

Not atomic, still requires multiple instructions, between which another process could be interleaved.

Locks

```

1 void Extract(int acc_no, int sum)
2 {
3     lock(L);
4     int B = Acc[acc_no];
5     Acc[acc_no] = B - sum;
6     unlock(L);
7 }
```



```

1 void lock(int *L)
2 {
3     while (L != 0)
4         /* wait */;
5     *L = 0;
6 }
```

```

1 void unlock(int *L)
2 {
3     *L = 0;
4 }
```

Here the implementation of the lock is flawed as we cannot atomically check the value of the lock, and set its value.

A working implementation is:

```

1 void lock(int *L)
2 {
```

```
3  while(TSL(L) != 0)
4      /* wait */ ;
5  }
```

TSL (test-and-set-lock) atomically sets the given memory location to 1 and returns the old value. Locks of this type are called **spin locks**.

Spin Locks

Spin locks should only be used when the expected wait time is short (otherwise wasteful). However it may run into the priority inversion problem.

Priority Inversion

When a lower priority thread is scheduled instead of a higher priority thread.

Threads $H > M > L$.

L acquires a lock, however is descheduled. H attempts to acquire the lock, but is blocked. The lock is not released as L needs to run in order to release it. M is higher priority than L , and runs instead.