

# 50004 - Operating Systems - Lecture 14

Oliver Killane

18/12/21

## Disk Scheduling

### (FCFS) First Come First Served


Requests completed in the order they were received.

- For light loads this is fine (time between requests is much larger than time taken to fulfil any request).
- Low performance for heavy loads (ends up traversing the disk more than optimal to get each request).
- Fair (no bias towards any process).

For example:

Head at 50. Queue: 195, 365, 75, 245, 30, 260, 125, 135

Operations: 50 → 195 → 365 → 75 → 245 → 30 → 260 → 125 → 135




### (SSTF) Shortest Seek Time First

Order requests according to shortest seek distance from the current head position

- Biased against innermost and outermost tracks (middle tracks on average closer, even with random head positions)
- Unpredictable performance (Requests may face a long delay if several very close requests are serviced)
- Processes can use dummy requests to keep control of head and have their requests prioritised.
- Can delay requests indefinitely (e.g many requests come in very close to each other, trapping the head).

Head at 50. Queue: 195, 365, 75, 245, 30, 260, 125, 135

Operations: 50 → 30 → 75 → 125 → 135 → 195 → 245 → 260 → 365




## SCAN Scheduling

Also called elevator scheduling. Select requests with the shortest seek time in the preferred direction.

- Only change direction when at the outermost/innermost cylinder (no more requests in the direction)
- Long delays for requests that are not in the path of the algorithm (e.g come in just behind the head) and for the extremes.
- Base for the most common algorithms used.
- Suffers from same delay issue as **SSTF**, though reduced (only in one direction)

Head at 160. Queue: 195, 365, 75, 245, 30, 260, 125, 135

Operations: 160 → 195 → 240 → 260 → 365 → 135 → 125 → 75 → 30




## C-SCAN Scheduling

Much like **SCAN** but scanning in one direction only, jumping to the start (e.g innermost) when at the end (e.g outermost)

- Lower variance of requests on extreme tracks
- Largely reduces issue of indefinite wait from **SSTF**.

Head at 160. Queue: 195, 365, 75, 245, 30, 260, 125, 135

Operations: 160 → 195 → 240 → 260 → 365 → 30 → 75 → 125 → 135




### N-Step SCAN Scheduling

Same as **SCAN**, however only services requests waiting when the sweep began (for each sweep)

- Requests arriving during sweep are serviced before end of the sweep (no long waits)
- No indefinite waits possible.

Head at 160. Queue: 195, 365, 75, 245, 30 Then 260, 125, 135

Operations: 160 → 195 → 245 → 365 → 75 → 30 → (*new sweep*) 125 → 135 → 260



## Linux Disk Scheduling

struct bio

Can be found here in the linux kernel.

```
1  /* main unit of I/O for the block layer and lower layers (ie drivers
2     ↪ and
3     * stacking drivers)
4     */
5  struct bio {
6      struct bio      *bi_next;    /* request queue link */
7      struct block_device *bi_bdev;
8      unsigned int    bi_opf;     /* bottom bits req flags, top
9     ↪ bits REQ_OP.
10
11     * Use accessors.
12     */
13     unsigned short    bi_flags;   /* BIO_* below */
14     unsigned short    bi_ioprio;
15     unsigned short    bi_write_hint;
16     blk_status_t      bi_status;
17     atomic_t          __bi_remaining;
18     struct bvec_iter    bi_iter;
19     bio_end_io_t      *bi_end_io;
20     void              *bi_private;
21
22 #ifdef CONFIG_BLK_CGROUP
23     /* Represents the association of the css and request-queue for
24     ↪ the bio.
25     * If a bio goes direct to device, it will not have a blkcg as
26     ↪ it will
27     * not have a request-queue associated with it. The reference
28     ↪ is put
29     * on release of the bio.
30     */
31     struct blkcg_gq      *bi_blkcg;
32     struct bio_issue      bi_issue;
33 #ifdef CONFIG_BLK_CGROUP_IOCOST
34     u64                  bi_iocost_cost;
35 #endif
36 #endif
37
38 #ifdef CONFIG_BLK_INLINE_ENCRYPTION
39     struct bio_crypt_ctx  *bi_crypt_context;
40 #endif
41
42     union {
43 #if defined(CONFIG_BLK_DEV_INTEGRITY)
44         struct bio_integrity_payload *bi_integrity; /* data
45         ↪ integrity */
46 #endif
47     };
48
49     unsigned short bi_vcnt; /* how many bio_vec's */
50
51     /* Everything starting with bi_max_vecs will be preserved by
52     ↪ bio_reset()
53     */
54     unsigned short bi_max_vecs; /* max bvl_vecs we can hold */
55     atomic_t      __bi_cnt; /* pin count */
56     struct bio_vec *bi_io_vec; /* the actual vec list */
57     struct bio_set *bi_pool;    4
58
59     /* We can inline a number of vecs at the end of the bio, to
60     ↪ avoid
61     * double allocations for a small number of bio_vecs. This
62     ↪ member
63     * MUST obviously be kept at the very end of the bio.
64     */
65     struct bio_vec    bi_inline_vecs[];
66 };
```

- **I/O Requests added to request list**

There is one request list per device, a **bio** that keeps track of the pages associated with the request.

- **Block Device drivers define a request operation for the kernel to use**

Interface provided by function pointers.

- Kernel passes ordered request list to driver.
- Driver completes all requests in the list.
- Drivers use the read/write operations defined by the kernel (uniform interface from kernel to send uniform request types to many different drivers).

- **Driver Based ordering**

Some drivers bypass kernel ordering and do it themselves (e.g **RAID**). This is done for more complex disk setups where assumptions made by the kernel ordering algorithm do not apply.

- **Default Algorithm: SCAN Scheduling**

Kernel attempts to merge requests to adjacent blocks (grouping adjacent requests results in less seek time, higher request throughput)

However read requests may starve during very large writes (from merging), if these are done synchronously, a program may hang.

- **Deadline Scheduler**

Ensures that reads are performed before a set deadline to prevent long or indefinite waits. (Eliminates read request starvation)

- **Anticipatory Scheduler**

Delay after read requests completed. If a process sends a second synchronous read request, it will be attended to quickly (due to delay after completing the first).

- Reduces Excessive seeking (Second request will likely be very close to the first, so avoid seeking away, and then back)
- Can reduce throughput (if no more read requests are sent by the process during the delay, or if the request sent requires a large seek)

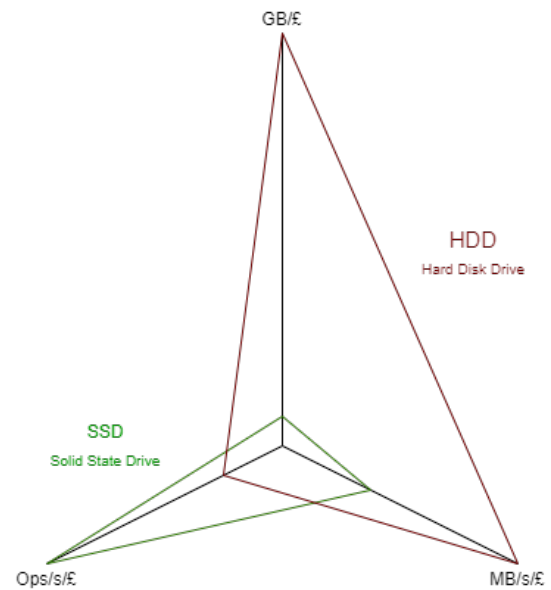
## Solid State Drives

### SSD Scheduling

Scheduling for SSDs do not require scheduling algorithms as many memory modules can be read/written in parallel and write/read speed is approximately constant.

However drivers need to overcome issues with limited writes, tracking virtual to physical blocks & assignment of free blocks.

- Very high bandwidth (e.g 1GB/s SSD vs 100MB/s HDD)
- Lower latency
- High parallelism



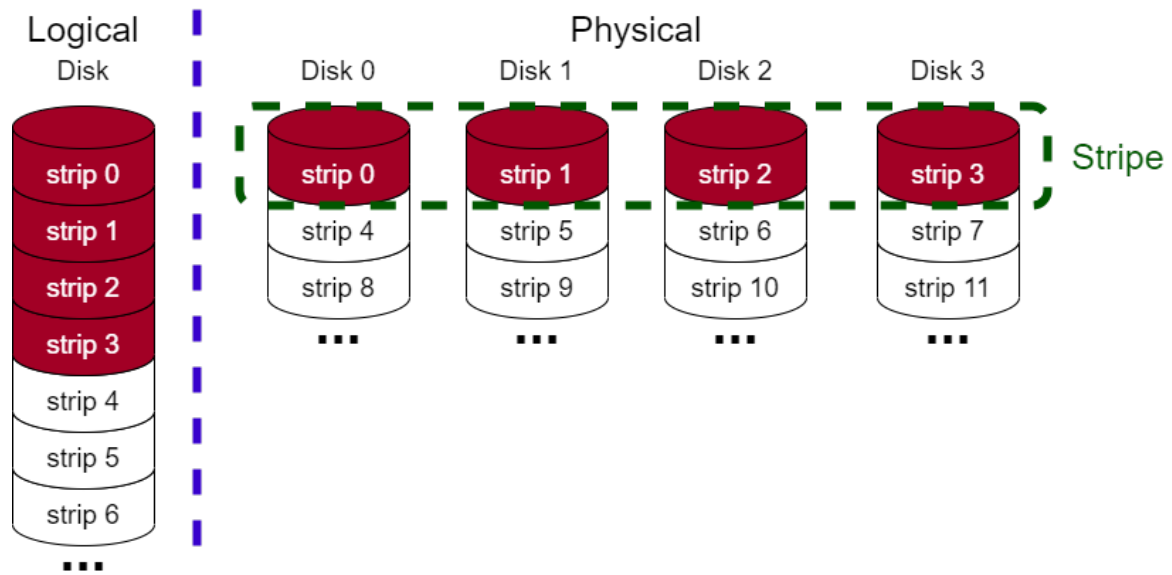
## RAID

Disk performance has not kept pace with CPU performance, creating a bottleneck. Redundant Array of Inexpensive Disks (**RAID**) increases disk based system performance by using many disks in parallel.

- An array of physical drives appears as a single virtual drive.
- Distributes stored data over physical disks to allow parallel operation, improving performance.
- Use redundant disk capacity to respond to disk failure (more disks means lower mean-time-to-failure, more disks, higher chance a single disk in the group fails).

There are several levels of **RAID**, each with differing performance, redundancy and space-efficiency/cost.

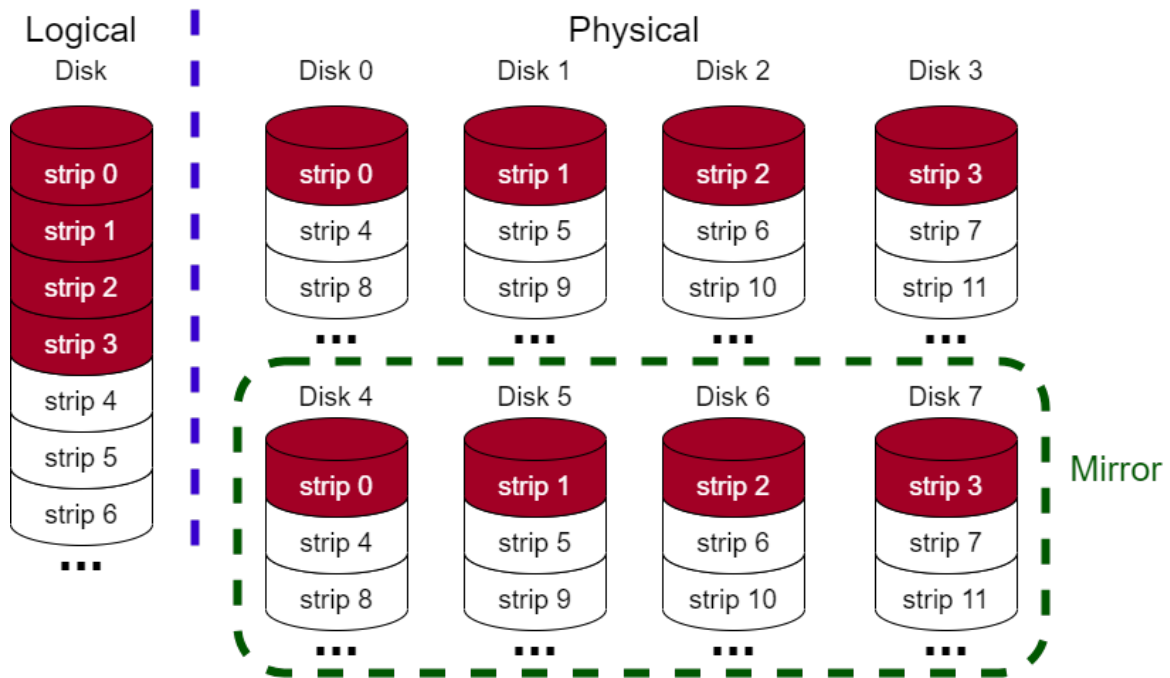
## RAID 0 - Striping



- Spread blocks in round robin fashion across disks.
- Can concurrently seek/transfer data (provided blocks on different physical disks).
- Can balance load across disks (sometimes).
- No redundancy (any disk failure will result in data being lost).



## RAID 1 - Mirroring



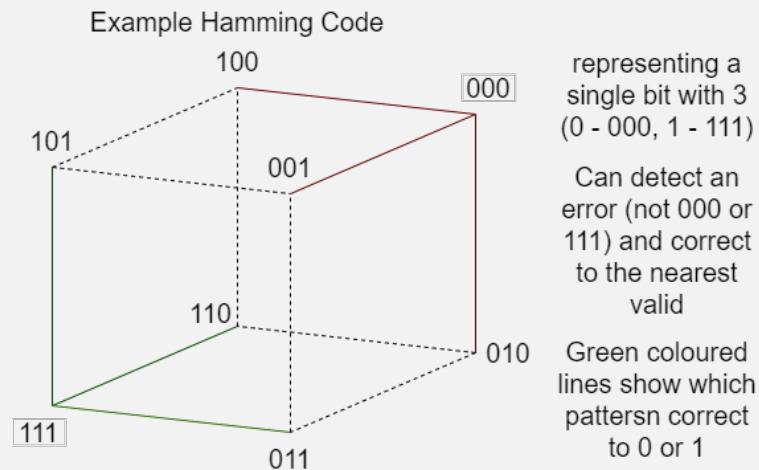
Mirror Data across disks to increase redundancy.

- Reads can be serviced by either disk (e.g can read 0, 1, 2, 3, 4, 5, 6, 7 at the same time).
- Writes must update both disks in parallel (e.g can effectively only write to 4 disks at a time). (Slower)
- Failure recovery is easy (when a disk fails, use its mirror).
- Low space efficiency and hence high cost (store everything twice).

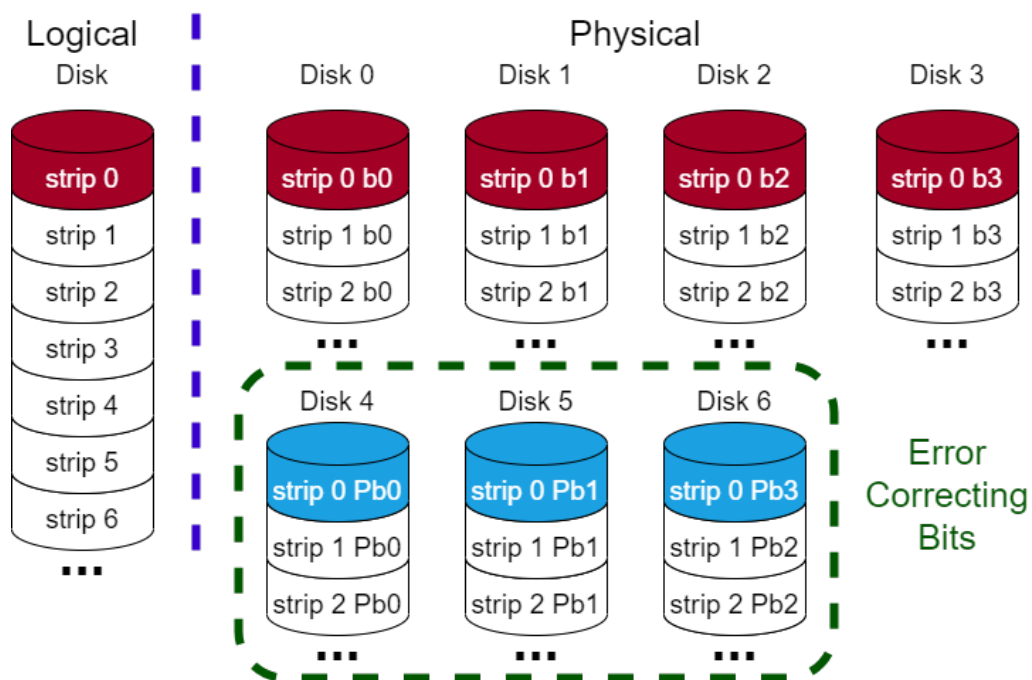
## RAID 2 - Bit-Level Hamming

### Hamming Codes

A family of error correcting codes that make use of hamming distance (number of bits different between two patterns) to correct single bit errors.



More complex codes can be used to check larger bit patterns with very few bits. The general algorithms can be seen [here](#)

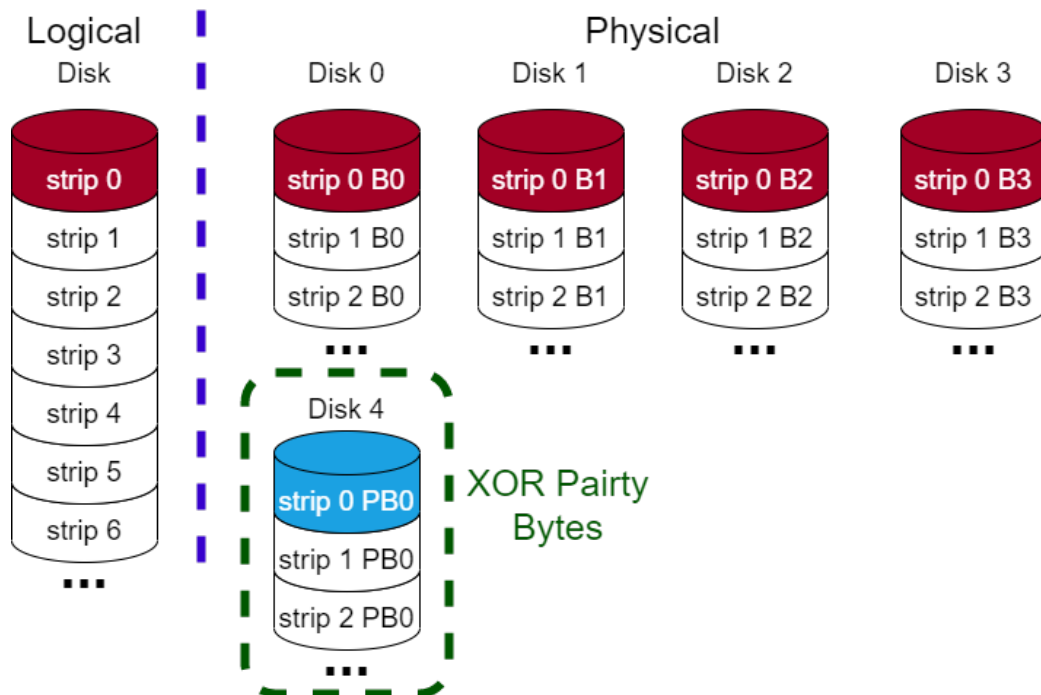


Parallel access by striping at the bit level.

- Consecutive (in this case 4) bits are read in parallel, hence very high throughput (always reading/writing in parallel).
- Hamming error-correcting codes used to detect and correct single bit errors (can detect but not correct double bit errors).
- Cannot process requests in parallel (each request requires all disks, no concurrency).
- High storage overhead (less space efficient, increasing cost).
- Large number of writes as the error correcting codes must be updated (can become a bottleneck) & every disk must write for every write operation.

## RAID 3 - Byte-Level XOR

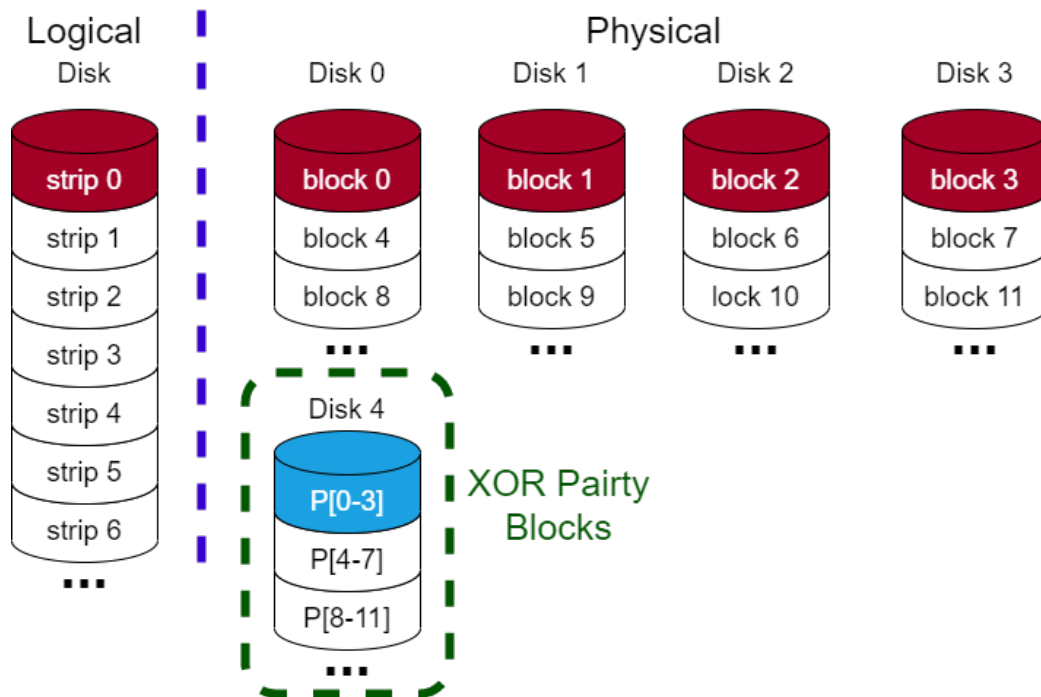
A single parity strip used (data on other disks XORed together).



Can reconstruct missing data from parity and remaining data when a disk fails (disk reports).

- Easy to reconstruct data when a disk fails
- More space efficient than **RAID 2**
- Only one I/O request at a time (but each request can be read in parallel from all disks)

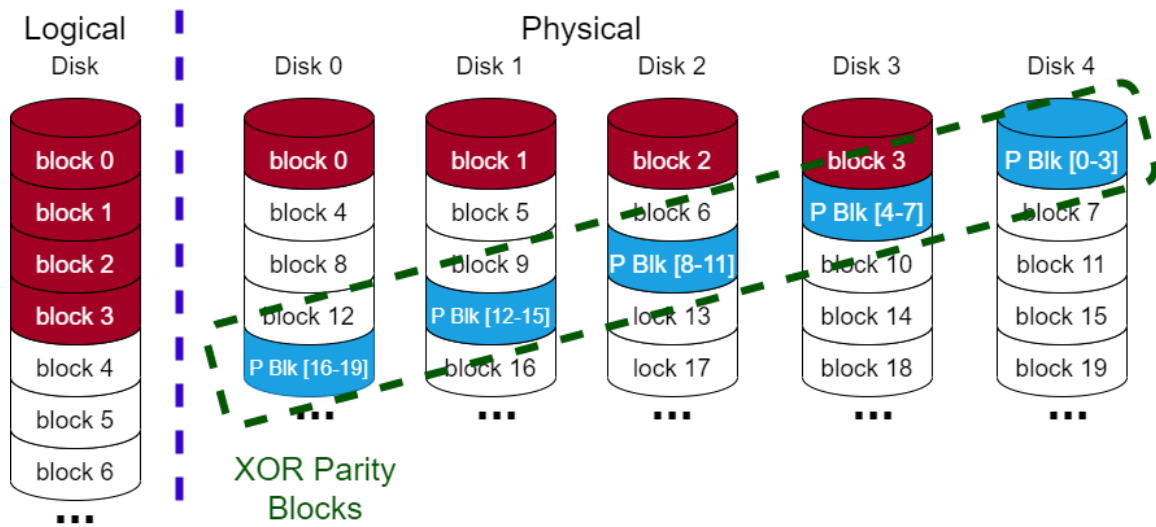
## RAID 4 - Block Level XOR



Parity Strip is XOR over blocks, allowing entire blocks to be accessed independently on different disks.

- Allows read requests to be serviced concurrently.
- Low redundancy overhead.
- Parity must be updated after every single write, Parity disk becomes a bottleneck.

## RAID 5 - Block Level Distributed XOR



The most commonly used **RAID** level, by distributing parity there is potential for concurrency.


- Some potential for write concurrency as parities are on different disks.
- Good storage/efficiency tradeoff.
- Reconstruction of disk is non-trivial (and slow).

## RAID Level summary

Speeds compared with using a single disk.

| Category           | Level | Description                          | I/O Data Transfer |       | I/O Request Rate |        |
|--------------------|-------|--------------------------------------|-------------------|-------|------------------|--------|
|                    |       |                                      | Read              | Write | Read             | Write  |
| Striping           | 0     | Non-Redundant                        | ↑                 | ↑     | ↑                | ↑      |
| Mirroring          | 1     | Mirrored                             | ↑                 | =     | ↑                | =      |
| Parallel Access    | 2     | Redundant (Hamming ECC)              | ↑↑                | ↑↑    | =                | =      |
|                    | 3     | Redundant (Bit Interleaved Parity)   | ↑↑                | ↑↑    | =                | =      |
| Independent Access | 4     | Block interleaved Parity             | ↑                 | ↓     | ↑                | ↓      |
|                    | 5     | Block interleaved distributed parity | ↑                 | ↓     | ↑                | = or ↓ |

## Disk Caching



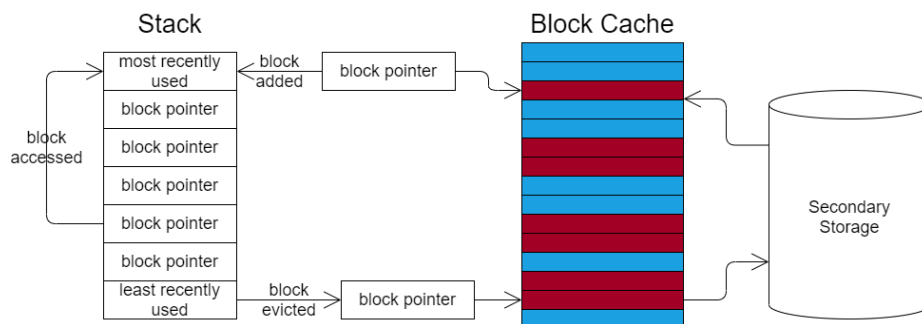
We can cache sectors of disk in main memory to reduce access times.

- Buffer contains copies of disk sectors.
- OS manages disk in terms of **blocks** which are likely much larger than sectors (so loads multiple sectors).
- Must ensure contents are saved in case of failure (e.g lazy writing is very complex).
- Cache has finite space, so a replacement policy must be implemented.

### (LRU) Least Recently Used

Replace the block that was in cache longest with no references.

Cache is a stack of pointers to blocks in memory, when a block is referenced, it is pushed to the top of the stack. Replacement evicts the block at the bottom of the stack.




This replacement policy does not track how many times a block is accessed, only the relative time of the last access (through stack position).

### (LFU) Least Frequently Used

- Replace the block with the fewest references.
- Each block has a counter (incremented when referenced), the lowest counter block is evicted during replacement.
- To avoid this, can use frequency based replacement.

## Frequency-Based Replacement



To prevent items that get a sudden burst of accesses from lingering for a long time in the cache, we only increment the count when a block is shuffled out of the **new** section.

## File Systems

File systems organise information. Their main objectives:

- **Non-volatile, long term storage**
- **Sharing information** e.g compilers, applications, text data etc.
- **Concurrent Access** Many programs can access many files pseudo-simultaneously
- **Convenient Organisation** symbolic names, directories
- **Easy management of data** automatic backup, snapshots etc.
- **Security** Permissions, read/write, hidden files

## File Naming

A **file** is an arbitrary size collection of data, extensions allow for easy identification of type from an identifier:

| Type        | Extension                   | Function  |
|-------------|-----------------------------|---|
| executable  | exe, com, bin, no extension | read & load to run machine code program.              |
| object      | obj, o                      | compiled machine language, not linked.                |
| source code | c, cpp, java, rs, py, s     | Source code, extension identifies language.           |
| batch/bash  | bat (windows), sh (unix)    | Script of commands to be interpreted by the terminal. |
| text        | txt                         | Text data (usually <b>ascii</b> ).                    |
| library     | lib, a, so, dll             | Libraries that programs can link to.                  |
| archive     | arc, zip, tar               | Compressed archives (for transmission or storage).    |

There are also several types of file.

| Type              | Description   |
|-------------------|---|
| Hard Links        | A file that aliases the data of another file (refers to the specific location of the data).   |
| Soft Links        | A soft/symbolic link aliases the path to a file by another (e.g file points to another file, which in turn points to its data). The <b>ln</b> command can be used to soft link files in unix based systems. |
| Regular           | Normal file as discussed in the table above.  |
| Directory         | A collection of files, that can itself be added to directories.   |
| Character Special | A file that provides access to a character I/O device   |
| Block Special     | Same as a character special file, but for block devices.  |

## Filesystem support functions

- **Name Translation** Converting paths to disks & blocks (logical) for use by the driver.
- **Management of Disk Space** Allocating and deallocating storage for files.
- **File locking for exclusive access** Important when many programs may access a file concurrently (can fuse the `fcntl` syscall with the `F_SETLK`, `F_SETLKW`, and `F_GETLK`)
- **Performance Optimisation** Caching/Buffering
- **Protection against system failure** Backup/Restore in case of a crash/power failure/unexpected shutdown
- **Security** Enforcing file permissions

## File Attributes

- **Basic**
  - **File Name** Symbolic name, unique in directory
  - **File Type** e.g text, binary, executable, directory
  - **File Organisation** placement of contents in blocks (sequential, random)
  - **File Creator** program that created the file
- **Address information**
  - **Volume** Disk drive, partition
  - **Start Address** cylinder, head, sector (logical block addressing information)
  - **Size Used**
  - **Size Allocated**
- **Access Control Information**
  - **Owner** User that controls the file
  - **Authentication** Locked files may need a password/key
  - **permitted actions** e.g read/write, delete permissions (owner/others)
- **Usage Information** Metadata for paper-trail.
  - **Creation Timestamp** Date & Time
  - **Last Modified** (can include the user that made the modification)
  - **Last Read**
  - **Last Archived** For keeping track of backups
  - **Expiry Date** For automatic deletion (e.g recycle bin contents)
  - **Access activity** Metadata for reads/writes etc. (can be used in improving performance)

## Unix/Linux Files

### Common Filesystem Syscalls

```
1  /* UNISTD functions (header for POSIX api) */
2  #include <unistd.h>
3  /* Close a file */
4  int close(int fd);
5
6  /* Read COUNT bytes from file FD into buffer BUF. Return number of bytes read. */
7  ssize_t read(int fd, void *buf, size_t count);
8
9  /* Write to file FD, COUNT bytes from BUF. Return number of bytes written. */
10 ssize_t write(int fd, const void *buf, size_t count);
11
```

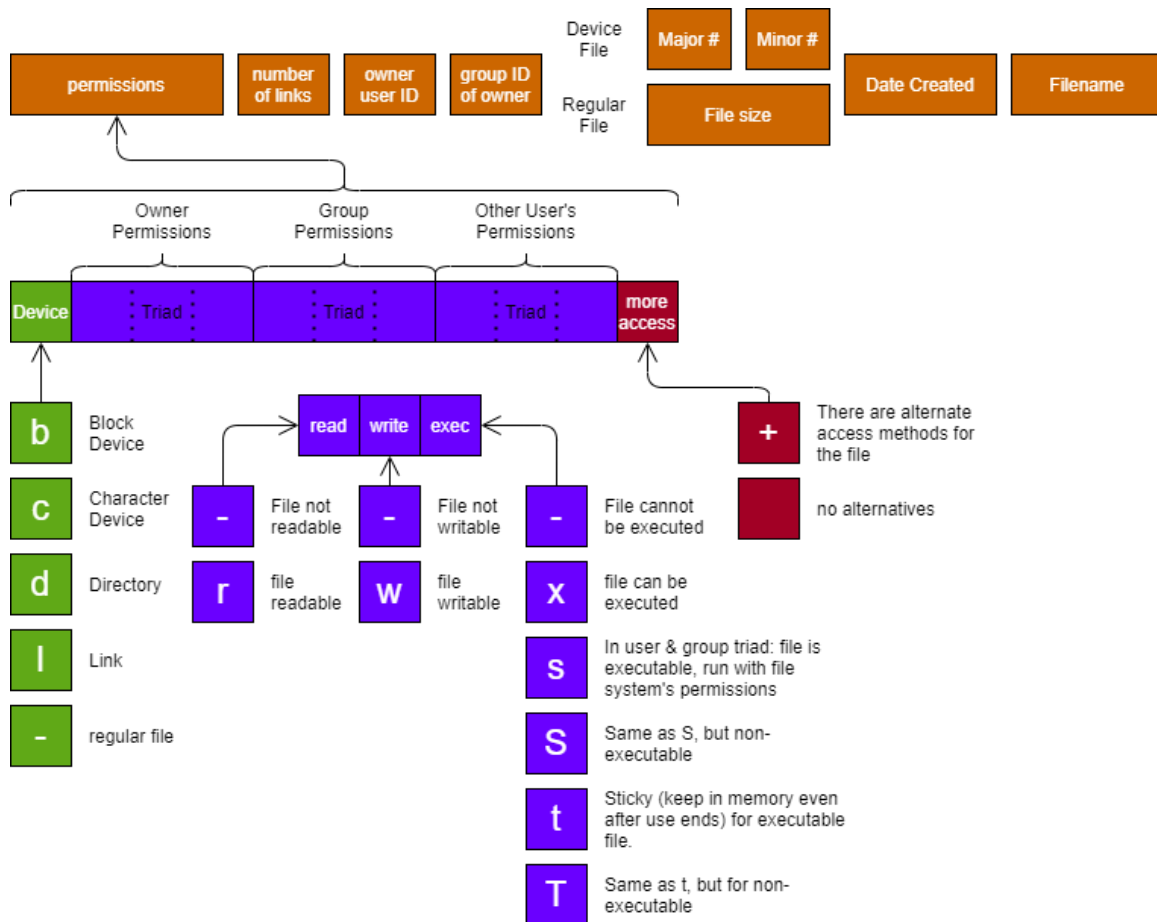


```

12 /* Seek to the given OFFSET in file FD, options in WHENCE. "lseek" stands for
13  * long seek (offset is 64 bits for large files).
14  */
15 off_t lseek(int fd, off_t offset, int whence);
16
17 /* Return file information for file at PATHNAME into buffer STATBUF. */
18 int stat(const char *restrict pathname, struct stat *restrict statbuf);
19
20 struct stat {
21     dev_t    st_dev;        /* ID of device containing file */
22     ino_t    st_ino;        /* Inode number */
23     mode_t   st_mode;       /* File type and mode */
24     nlink_t  st_nlink;      /* Number of hard links */
25     uid_t    st_uid;        /* User ID of owner */
26     gid_t    st_gid;        /* Group ID of owner */
27     dev_t    st_rdev;       /* Device ID (if special file) */
28     off_t    st_size;       /* Total size, in bytes */
29     blksize_t st_blksize;    /* Block size for filesystem I/O */
30     blkcnt_t st_blocks;     /* Number of 512B blocks allocated */
31
32     /* Since Linux 2.6, the kernel supports nanosecond
33      * precision for the following timestamp fields.
34      * For the details before Linux 2.6, see NOTES.
35      */
36
37     struct timespec st_atim; /* Time of last access */
38     struct timespec st_mtim; /* Time of last modification */
39     struct timespec st_ctim; /* Time of last status change */
40
41     #define st_atime st_atim.tv_sec      /* Backward compatibility */
42     #define st_mtime st_mtim.tv_sec
43     #define st_ctime st_ctim.tv_sec
44 };
45
46
47 /* File Control functions */
48 #include <fcntl.h>
49
50 /* Note that interestingly, despite close being unistd, open is not. This is as
51  * many of the flags required for open are declared in FCNTL, and the developers
52  * wanted to avoid polluting UNISTD with these.
53  */
54
55 /* Open a file at PATHNAME with FLAGS to get a file descriptor. */
56 int open(const char *pathname, int flags);
57 int open(const char *pathname, int flags, mode_t mode);
58
59 /* File controls (e.g lock, set read status). */
60 int fcntl(int fd, int cmd, ... /* arg */ );

```

## File Attributes



```

1 nolinks# For a device:
2 # <permission attributes> <no. links> <owner> <group> <major> <minor> <creation date>
  ↪ <filename>
3
4 crw----- 1 root root    5,   1 Dec 20 14:59 console
5 # Character device, readable and writeable, but not executable by the owner
6
7 brw----- 1 root root    1, 15 Dec 20 14:59 ram15
8 # Block device, readable and writeable, but not executable by the owner
9
10 # For a file:
11 # <permission attributes> <no. links> <owner> <group> <size> <creation date> <
  ↪ filename>
12
13 -rwxrwxrwx 1 oliverkillane oliverkillane 21775 Dec 20 16:31 'Lecture 14.tex'
14 # File, readable, writable and can be executed by any user

```

# File System Organisation

## Space Allocation

File size is variable (can increase or decrease) and is allocated on the disk in blocks (usually 512 → 8192 bytes). The block size is determined by the file system.

### Block Size too Small

- High overhead for managing large files (large number of blocks to keep track of)
- High file transfer time (large file → lots of blocks across the disk → lots of seeking back and forth)


### Block Size too Large

- Internal fragmentation (Files leave parts of large blocks unused - wasteful)
- Small files waste lots of space.
- Caching is based on blocks, so end up having a very large cache space, or unable to cache many blocks.

## Contiguous File Allocation

Place file data on a contiguous stretch of addresses on storage device. Similar to segment based memory management.

- Successive logical records are usually physically adjacent, reducing seek time when reading the file (seek to start, then just traverse file).
- External Fragmentation can occur (unused blocks between files).
- If unable to allocate new blocks as a file grows, must transfer to new large free section (expensive, requires lots of I/O).




## Block Linkage/Chaining

Each file contains a linked list of blocks. When locating a block, traverse the linked list to by seeking to block, reading pointer.

- Can grow files without expensive re-allocation (just set pointer of end block to point to a newly allocated one).
- No external fragmentation (as all blocks can be used in any order, by any file).
- Space used for pointer in each block.
- Traversing files requires lots of seeking from block to block.

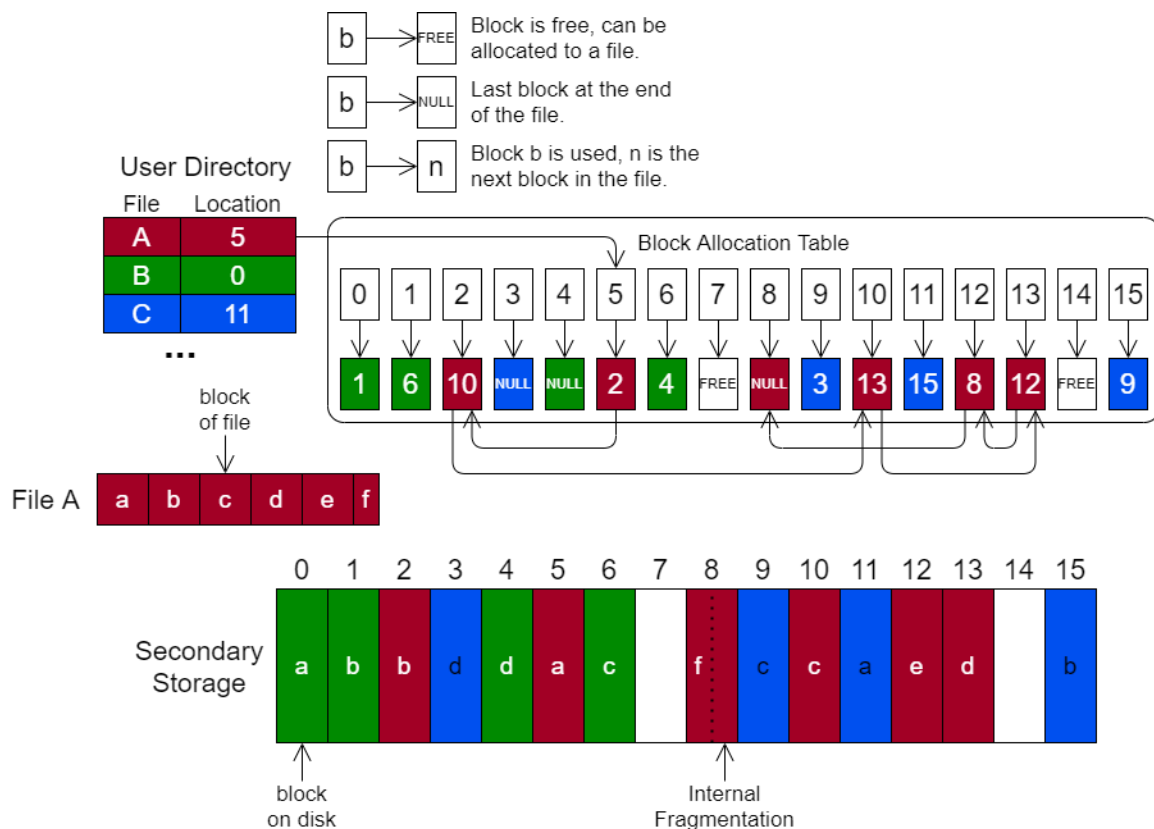
- For small block sizes the number of seeks to traverse a file increases.



## Block Allocation Table

Uses a directory mapping files to first block. Table maps blocks to the next block in the file, indicating free spots (**FREE**) and block with no next (end of file - **NULL**).

- File allocation table can be cached in memory for fast lookup.
- Does not require lengthy seeks to traverse the block numbers for the file.
- No external fragmentation.
- However files can become fragmented (spread across disk) which reduces read/write speed (e.g must do lots of seeks to read the whole file), so should be periodically defragmented (disk reorganised to place files blocks next to each other, very expensive operation).
- Table can become impractically large, using up lots of memory, and more than one block of storage.

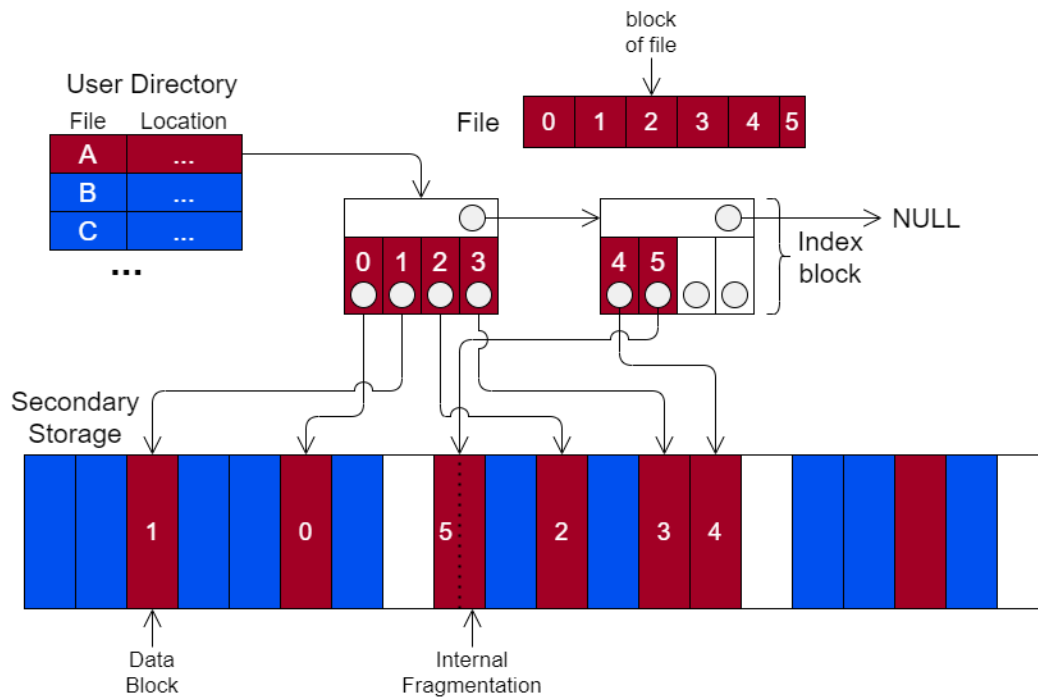


This system is used for **microsoft's FAT16/32** file system (with the table cached in memory).

## Index Blocks

Each table has one or more index blocks. Index blocks contain an array of pointers to data blocks, and pointers to subsequent index blocks. Basically page tables for file systems. The example below has miniscule index blocks for illustration purposes.

- Can search through index blocks to find location fo file blocks easily.
- No external fragmentation & can extend files easily.
- Index blocks can be cached in memory just like data blocks.
- Index blocks are per-file, hence can load file's table, rather than have an enormous global table as with **FAT**



## Linux Inodes

Linux/Unix uses the **Index Blocks** strategy through **inodes** (index nodes).

Inode contain:

- Type & Access control.
- Number of links to that inode.
- User & Group ID.
- Access & Modification time.
- Inode change time (e.g when permissions, data stored in inode were changed).
- Direct, Indirect, Double Indirect & Triple Indirect pointers to data blocks.

The **inode** struct can be found here in the linux kernel.

```

1 struct inode {
2     umode_t      i_mode;
3     unsigned short i_opflags;
4     kuid_t       i_uid;
5     kgid_t       i_gid;
6     unsigned int  i_flags;
7
8 #ifdef CONFIG_FS_POSIX_ACL
9     struct posix_acl *i_acl;
10    struct posix_acl *i_default_acl;
11 #endif
12
13    const struct inode_operations *i_op;
14    struct super_block *i_sb;
15    struct address_space *i_mapping;

```

```

16
17 #ifdef CONFIG_SECURITY
18     void                *i_security;
19 #endif
20
21     /* Stat data, not accessed from path walking */
22     unsigned long        i_ino;
23     /*
24      * Filesystems may only read i_nlink directly. They shall use the
25      * following functions for modification:
26      *
27      * (set|clear|inc|drop)_nlink
28      * inode_(inc|dec)_link_count
29      */
30     union {
31         const unsigned int i_nlink;
32         unsigned int __i_nlink;
33     };
34     dev_t                i_rdev;
35     loff_t                i_size;
36     struct timespec64     i_atime;
37     struct timespec64     i_mtime;
38     struct timespec64     i_ctime;
39     spinlock_t            i_lock; /* i_blocks, i_bytes, maybe i_size */
40     unsigned short        i_bytes;
41     u8                    i_blkbits;
42     u8                    i_write_hint;
43     blkcnt_t              i_blocks;
44
45 #ifdef __NEED_I_SIZE_ORDERED
46     seqcount_t            i_size_seqcount;
47 #endif
48
49     /* Misc */
50     unsigned long         i_state;
51     struct rw_semaphore   i_rwsem;
52
53     unsigned long         dirtied_when; /* jiffies of first dirtying */
54     unsigned long         dirtied_time_when;
55
56     struct hlist_node     i_hash;
57     struct list_head      i_io_list; /* backing dev IO list */
58 #ifdef CONFIG_CGROUP_WRITEBACK
59     struct bdi_writeback  *i_wb; /* the associated cgroup wb */
60
61     /* foreign inode detection, see wbc_detach_inode() */
62     int                    i_wb_frn_winner;
63     u16                    i_wb_frn_avg_time;
64     u16                    i_wb_frn_history;
65 #endif
66     struct list_head      i_lru; /* inode LRU list */
67     struct list_head      i_sb_list;
68     struct list_head      i_wb_list; /* backing dev writeback list */
69     union {
70         struct hlist_head i_dentry;
71         struct rcu_head i_rcu;
72     };
73     atomic64_t            i_version;
74     atomic64_t            i_sequence; /* see futex */
75     atomic_t              i_count;

```

```

76     atomic_t    i_dio_count;
77     atomic_t    i_writecount;
78 #if defined(CONFIG_IMA) || defined(CONFIG_FILE_LOCKING)
79     atomic_t    i_readcount; /* struct files open RO */
80 #endif
81     union {
82         const struct file_operations    *i_fop; /* former ->i_op->
83             ↪ default_file_ops */
84         void (*free_inode)(struct inode *);
85     };
86     struct file_lock_context *i_flctx;
87     struct address_space    i_data;
88     struct list_head        i_devices;
89     union {
90         struct pipe_inode_info *i_pipe;
91         struct cdev            *i_cdev;
92         char                    *i_link;
93         unsigned               i_dir_seq;
94     };
95     __u32 i_generation;
96
97 #ifdef CONFIG_FSNOTIFY
98     __u32 i_fsnotify_mask; /* all events this inode cares about
99     ↪ */
100     struct fsnotify_mark_connector __rcu    *i_fsnotify_marks;
101 #endif
102 #ifdef CONFIG_FS_ENCRYPTION
103     struct fscrypt_info    *i_crypt_info;
104 #endif
105
106 #ifdef CONFIG_FS_VERITY
107     struct fsverity_info    *i_verity_info;
108 #endif
109
110     void *i_private; /* fs or device private pointer */
111 } __randomize_layout;

```