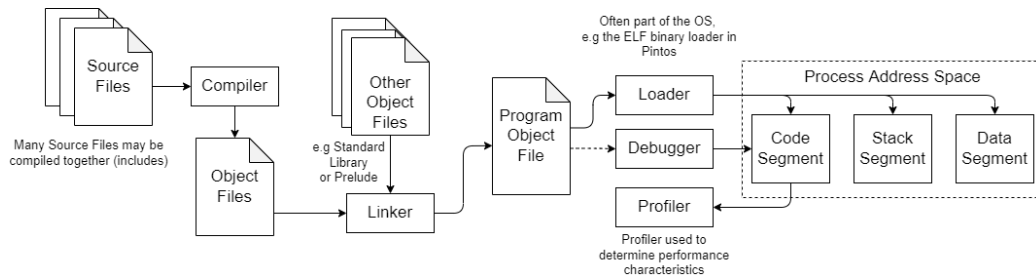


50006 - Compilers - (Dr Dulay) Lecture 5

Oliver Killane

14/04/22

Classic Approach to Compilers



Data Representation

Definition: Primitive Types

The most basic types, supported by the architecture the compiler targets.

Type	Bytes	Representation
Boolean	1	0 for false and 1 for true.
Integer	1,2,4,8	Typically 2s-complement.
Unsigned Integer	1,2,4,8	Basic unsigned binary.
Float/Real	4,8,16	IEEE Floating Point.
Char	1,2	For ascii (byte) and unicode.

- Compilers may align types for more optimal memory access, sometimes this is enforced by the target architecture.
- Some languages support pointer types (e.g C), however this makes type checking difficult (consider pointer arithmetic). As a result more modern language use type-checkable object references.

Definition: Alignment

Typically some address is considered n -byte aligned when it is a multiple of n bytes.

- Often data is aligned by the size of data the processor can load in a single instruction. This is done to ensure loading some data requires the minimal number of memory accesses.
- Some architectures enforce alignment for this reason, particularly for the stack.
- When aligned to a multiple of two 2^k , the first k bits of the address will be zero. For example alignment of pages in memory allow for page table entries to use the bits that will always be zero (due to alignment) to be used for other information (e.g reference, supervisor, read/write)

Definition: Structs/Records

Data types consisting of groups of fields/members of other types.

- Fields typically allocated in a contiguous block of memory.
- Alignment often used to space fields.
- Some languages order fields by their position in code (e.g C) while others can reorder and optimise.
- Tuples are effectively anonymous structs.

If no padding/alignment is used, and the struct/record is kept in order as expressed by the source code (e.g typical of C):

$$Record = (Field_1, Field_2, \dots, Field_n)$$

$$Size(Record) = Size(Field_1) + Size(Field_2) + \dots + Size(Field_n)$$

$$Address(Field_k) = StartAddress(Record) + Size(Field_1) + \dots + Size(Field_{k-1})$$

Struct Representation

As previously mentioned some languages reorder and optimise structs/records.

Rust not only does this, but allows the programmer to specify different representation schemes through macros. (See Rust alternative representations.)

Definition: Arrays

A contiguous section of memory populated by n variables (elements) of the same type.

- Can have elements aligned
- Some languages associated arrays with auxiliary data (e.g length for bounds checking).

A basic array implementation (elements not aligned/padded):

$$\text{Array}[\text{Type}] = \text{Element}[\text{Type}]_1, \dots, \text{Element}[\text{Type}]_n \text{ (Indexed as 0 to } n - 1)$$

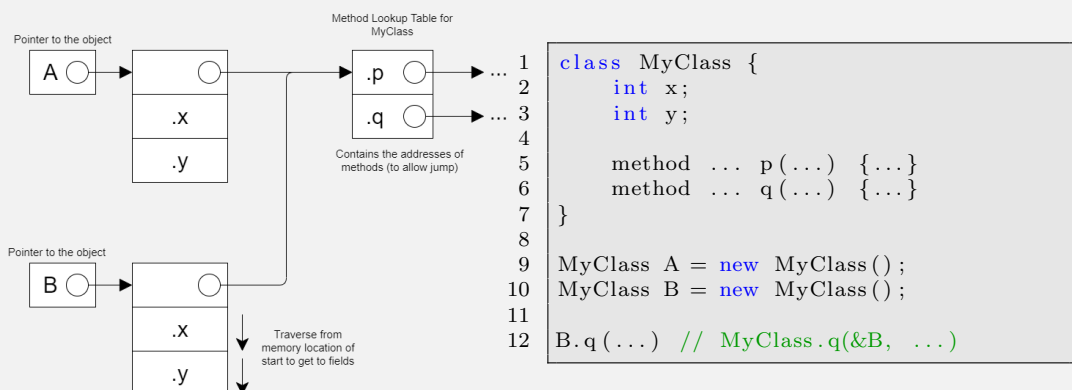
$$\text{Size}(\text{Array}) = \text{Size}(\text{Type}) \times n$$

$$\text{Address}(\text{Element}_k) = \text{StartAddress}(\text{Array}) + k \times \text{Size}(\text{Type})$$

Definition: Objects

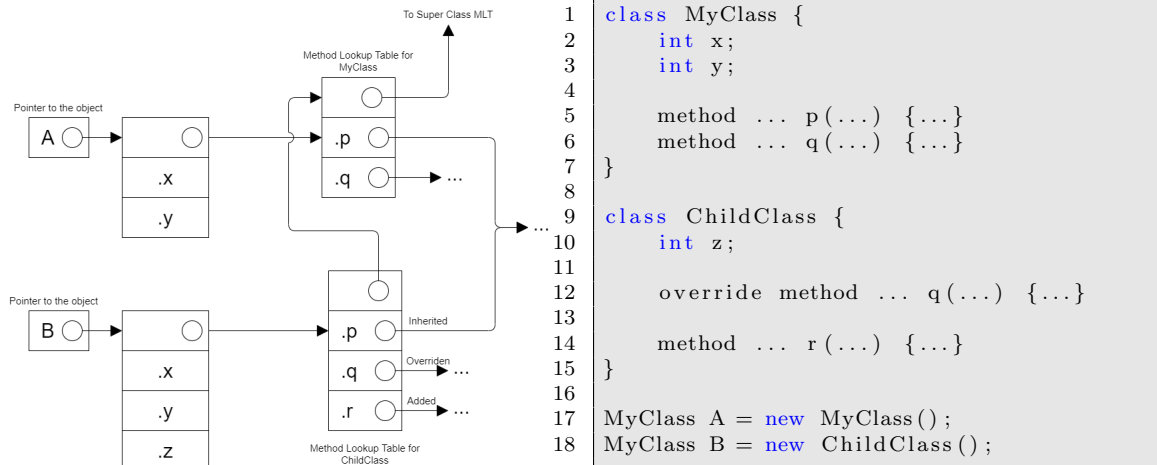
Implemented as a reference to a record, but with a pointer to a **method lookup table(MLT)** for that class (needs to consider inheritance).

- When calling a method of an object, traverse to the object's **MLT**. Then jump to the method, having placed the first argument as a pointer to the object.
- Access fields just like a struct.
- The indirection required for accessing values, calling methods adds overhead.



Class Inheritance

For single inheritance we include the parent's fields and methods in the struct and **MLT**.



- Pointers in the **MLT** go to the method (potentially a method also pointed to by the parent if no overridden)
- New added methods are in the **MLT**
- We chain the child/subclass's **MLT** to the parent/superclass' so that we can check for types (e.g. is this an *instance of* another class, a child etc) at runtime using the **MLT**. We could also have used type descriptors or other implementations.
- We preserve the layout of the parent/super class in the child (only appending methods and fields).

Dynamic Binding

As we only append fields and entries to the **MLT** when inheriting, it is possible to set a variable of type parent class to a value of type child class.

In the example below, we can see that a new field is added in *B*, (the inherited *A.n* is still available). The **MLT** of the *B* class contains the overridden *get* function.

```

1  public class OOPEXample {
2      public static void main(String args[]) {
3          A a = new B();
4          // As get returns the class's n, and the get method is B's we get 2.
5          // However a.n refers to A.n, which is 1. The field n is inherited from
6          // A, B adds another field called n.
7
8          // Outputs: "a = new B(); a.get() -> 2, a.n -> 1"
9          System.out.printf("A a = new B(); a.get() -> %d, a.n -> %d", a.get(), a.n);
10     }
11 }
12
13 /**
14  * A:
15  * -> MLT: A.get (defined)
16  * - int A.n (defined)
17  */

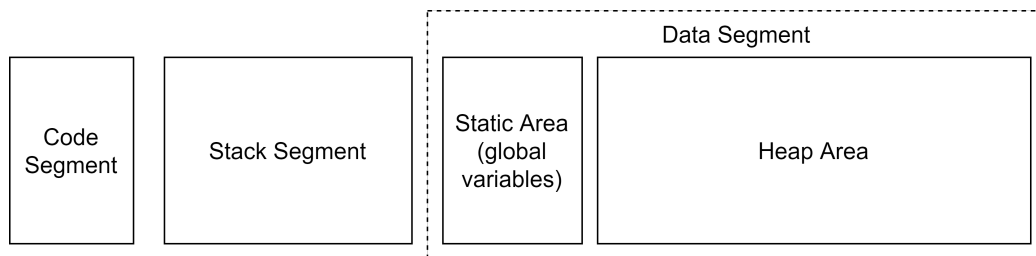
```

```

18 class A {
19     public int n = 1;
20     public int get() { return n; }
21 }
22
23 /**
24  * A:
25  * -> MLT: B.get (override)
26  * - int A.n      (inherited)
27  * - int B.n      (added)
28  */
29 class B extends A {
30     public int n = 2;
31     public int get() { return n; }
32 }

```

Program Address Space



Typically programs divide their address space into several segments (each potentially multiple pages large).

- By placing static data, code and stack space in different pages (in different segments/parts of the address space) they can have different permissions (e.g code is read-execute only, stack cannot be executed, etc)
- Global variables can be placed in a fixed location in the static area.
- Constants, **MLTs** and other global data can be kept in the static area.
- Local variables are stored on the stack in the stack frame of the function they are local to. They are allocated for the duration of the call.
- The stack can also contain arguments sent to the function, as well as information such as the address of the object (for methods) and slots for temporarily storing registers.
- Some architectures such as IA32 have a special register for the start address of the current stack frame (in IA32 it is the EBP register).

Heap Management

Lecture Recording

Lecture recording is available here

We can heap allocate variables (heap variables, also called dynamic variables). Here they can remain for any period of time, and be referenced from anywhere within the program (unlike local variables which are bound to a given function).

Heap variables must be allocated (e.g by *new* in java, or *malloc* in C/C++)

Programming languages manage memory in 4 different ways:

Explicit	C, ASM	Programmer determines when memory is freed, and how it is accessed. This is difficult, but allows for very fast code as the programmer can highly optimize their memory usage.
Garbage Collection	Java, Haskell, Go	The compiler adds code to check what objects are no longer in use, and to free them. This is done at runtime and has considerable overhead.

Rust Superiority

Rust's main memory management tool is Lifetimes (though it does use smart pointers with reference counting, and allow for explicit memory management also).

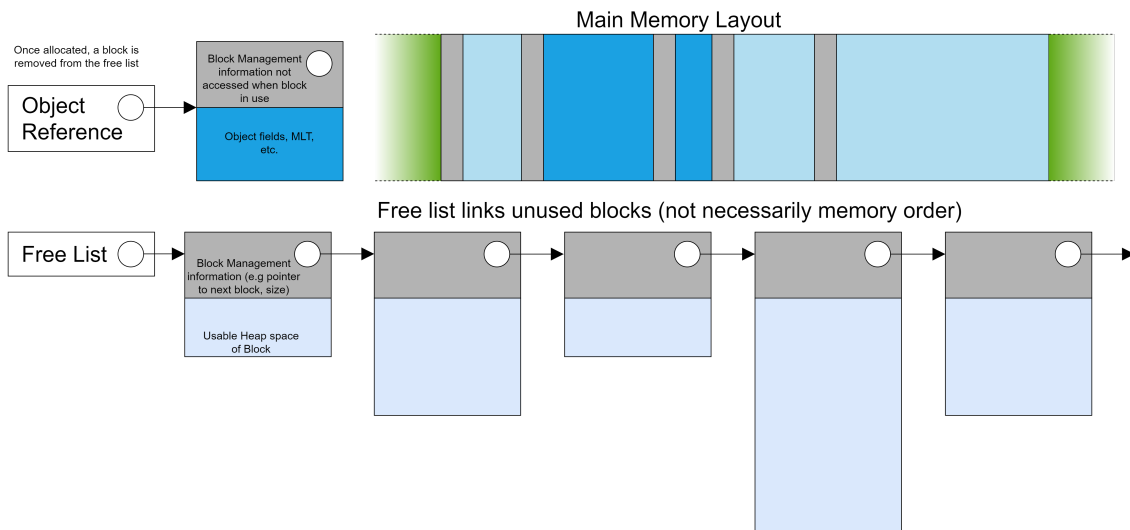
All references are qualified with information on mutability, and where the reference will be valid. The compiler uses these to ensure memory safety and insert required frees at compile time.

This makes the compiler more complex, but is safer than a GC (fewer bugs, will refuse to compile unsafe code (e.g with data races)), with performance almost identical to explicit memory management.

Explicit Heap Allocation

We maintain a structure containing the free blocks and information associated with each block. Here we will consider a simple linked list.

- Maintaining several lists for different sizes can reduce the search time.
- Should allocate blocks just large enough.
- For extensible arrays (vectors, mutable strings) we should allocate with the expectation that more memory will be used (allocate more than currently required).
- Some architectures require alignment of values, this must be considered. Requirements aside, as previously mentioned not-aligning can reduce performance.
- Heap allocation is worse for locality (e.g than the stack) and hence can potentially have worse cache locality, affecting performance.



```

1  def allocate(size: int) -> Address:
2      # Search through the blocks, can use a better algorithm:
3      block = free_blocks.search(size)
4      if block.size == size:
5          # found exactly the size needed
6          free_blocks.remove(block)
7          return block.start_address
8      elif block.size > size:
9          # split the block
10         (size_block, other) = free_blocks.split(size)
11         free_blocks.remove(size_block)
12         return block.start_address
13     else:
14         # we need more memory
15         if os.give_me_more_mem():
16             # try again with allocation
17         else:
18             # OS could not provide
19             return NULL
20
21 # Can be implemented to takes the block size as a parameter also
22 def deallocate(ptr: Address):
23     # Check free is valid
24     if free_blocks.check(address):
25         # return block to free blocks, coalesce/merge with another block to
26         # reduce fragmentation
27         free_blocks.insert_and_coalesce(ptr)

```


Garbage Collection

Definition: Garbage Collector

A part of the language's runtime that ensures heap allocated variables are properly de-allocated when they are no longer used. Many different **GC** algorithms are employed by many different languages.

GCs have several requirements:


- **Correctness** The **GC** must never collect (de-allocate) live data (in use).
- **Performance** Must be fast & have low memory overhead to reduce the impact on program performance.
- **Compiler Supported** In order to function the compiler needs to provide the garbage collector information for:
 - Which variables point to the heap.
 - The pointers contained in each block (e.g an object contains a reference to an object).

The compiler can either provide this data directly (type description data for objects), or generate special subroutines for the **GC** to call.

Garbage collectors can be run as **1-shot** (pause the program, run the **GC** and then resume), **on-the-fly** (as the program is run) or concurrent (runs on a different thread).

Definition: Heap Compaction

After **GC** de-allocates many blocks, the heap may be fragmented (lots of small free blocks). Hence much memory may not be free, but it may not be possible to allocate for an object as there is no block large enough.



Definition: Reference-Counting Garbage Collector

The block's management/housekeeping information contains a **reference count**. When a reference is made to an object, the reference count of the block containing it is increased, when that reference is removed, this is decreased.

When the reference count is decremented to zero, we know to free the block.

- Simple and efficient (garbage collection done as the program runs)
- Requires the compiler to generate the code to correctly track references.
- In reference cycles (e.g A references B which references A, even when no-one references them (they should both be collected) they keep each other alive) blocks may be kept live indefinitely, the **GC** needs to avoid this.

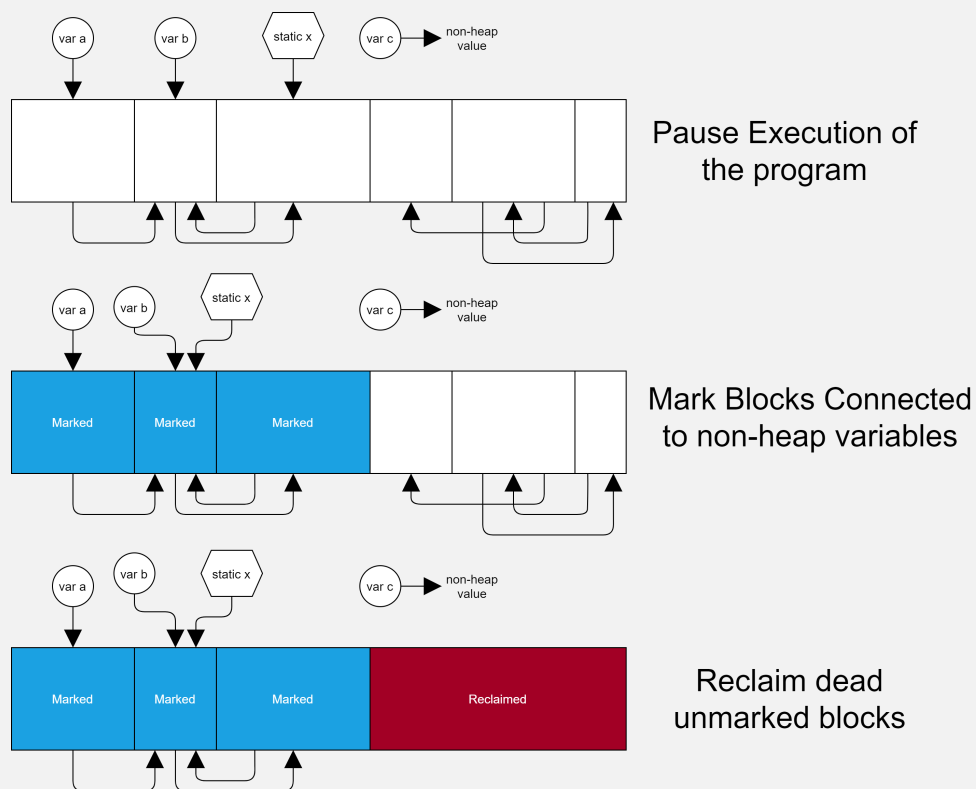
```
1 # used for assignments such as a = b
2 def assign(lval, rval):
3     if rval.on_heap():
4         rval.references += 1
5
6     if lval.on_heap():
7         lval.references -= 1
8         if lval.references == 0:
9             gc.reclaim(lval.block)
10
11     lval.value = rval.value
```

Definition: Mark-Sweep Garbage Collector

Pause the program, and collect all blocks that are not pointed to (potentially indirectly) by a non-heap value (e.g. local or static variable).

This will find all dead blocks/garbage (unlike a **reference counting GC**) however it has considerable overhead. Each block's maintenance/housekeeping information must also contain space for a mark bit.

- **Mark** For all stack variables, if they point to the heap then recursively traverse, marking each block they point to as live.
- **Sweep** Go through each block, if it is not marked as live, then collect and deallocate it.



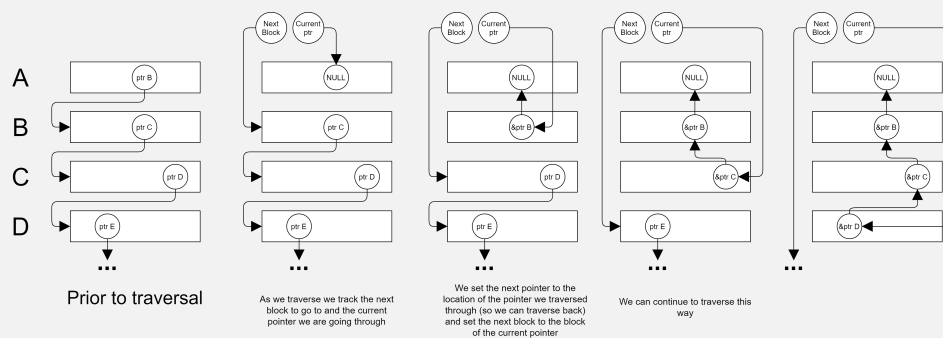
Definition: Pointer-Reversal Marking

Recursive traversal of pointers requires a potentially large amount of stack space (e.g going through a linked list of thousands of elements).

A garbage collector needs to work when memory is low. Hence we can use **Pointer-Reversal Marking** to store the pointers to return to (previous in traversal) in place.

1. When visiting a block, we select a pointer from within the block. We then set that pointer as the current, and place the address of the previous current pointer into it (hence now reversed).
2. We can then use the current pointer value to traverse to the next block, and continue.
3. When we find an already marked block, or no pointer in the block to traverse to, we can retreat back by using the pointer held in the location of the "current pointer" as this points back to the parent. We can set this
4. Note we need to store the number of marked children in the block's maintenance/-housekeeping data.

By storing the previously visited in the pointer we are traversing through, we require no extra stack space.



```

1 // Variables used to keep track of location
2 void **current_ptr;
3 void **next_ptr;
4 void* current_block;
5
6 // traverse downwards
7 next_ptr = find_ptr_in_block(current_block);
8
9 current_block = *next_ptr; // current_block = address of next block
10 *next_ptr = current_ptr; // next_ptr now contains the address of where
    ↳ the old current_ptr was stored
11 current_ptr = next_ptr;
12
13 // we can mark the block we are current at
14 mark_block(current_block);
15
16 // traversing upwards (assuming no next_ptr)
17 next_ptr = *current_ptr; // the next pointer is now previous (
    ↳ traversing back)
18 *current_ptr = current_block; // set the current ptr back to being the
    ↳ next block
19 current_ptr = next_ptr; // current pointer moves to the previous
20 current_block = block_from_ptr(*current_ptr); //get the block we are
    ↳ now in

```

Definition: Two Space Garbage Collector

The heap is split into two, a **from-space** and a **to-space**. Blocks are allocated from the **from-space**. When there are no more free blocks, all live (reachable from a non-heap pointer) blocks are copied to the **to-space** and then the **to-space** becomes the **from-space** and vice-versa.

- Fast, there are few complex pointer manipulations.
- Automatically compacts memory when copying.
- Can place linked blocks close together in memory when copying to allow for better cache locality.

Definition: Generational Garbage Collector

The heap is split into several areas based on the age of blocks.

- Allocate new blocks from the youngest, if full, move the oldest young blocks to an older area.
- **GC** is used more frequently on the younger areas.
- Can apply different **GC** techniques to the different areas.

Java 13

Java 13 has 5 different garbage collectors, you can read more about there here.

- Serial
- Parallel
- Concurrent Mark-Sweep (CMS)
- Garbage First (G1),
- ZGC

Other Compiler Components

Definition: Debugger

Useful tools for inspecting the behaviour of programs.

- Program behaviour must be maintained when being debugged (execution time can change, but correctness must be consistent).
- Debugging information (e.g function and variable names, source line mappings) may be embedded in the program.
- Highly optimising compilers and debuggers are mostly incompatible (as a useful debugger can map behaviour back to source code, and no such relation exists once complex optimisations are done). Hence debuggers usually compile with optimisations off.

There are two main kinds of debuggers:

- **Interactive** e.g GDB Allow users to inspect the state of the program, variables, functions, memory etc and to pause execution of the program.

For natively compiled (e.g Rust, C, Go) the architecture and OS must support breakpoints. For bytecode compiled (e.g Java) the interpreter supports this.

- **Post-Mortem / Core Dump** e.g Rust backtrace Provides debugging information upon failure by doing a reverse lookup from the program counter.
 - Stack TraceBack shows the methods called, their local variables, arguments etc.
 - Contents of global and dynamic/heap variables.

This requires the debugger to work out where variables are stored (register, stack, heap) and to map them back to names from the source using debugging information embedded at compile time.

Definition: Profilers

Provide performance information on a program.

- Typically used once code is correct, algorithmically optimal and optimised by the compiler.
- Can determine the time spent in functions, or how much of a program is spent in certain sections of code.

For example a profiler could interrupt execution periodically (some number of ms), identify the method being used, and increment a counter for it. Hence after profiling we can get a breakdown of the proportion of program time spent in each function.